

 h2oai / h2o-tutorials

 Watch 144

 Star 232

 Fork 229

<> Code

Issues 5

Pull requests 2

Projects 0

Pulse

Graphs

Branch: master ▾ h2o-tutorials / tutorials / deeplearning / deeplearning.R

Find file

Copy path

 angela0xdata PUBEV-3632 Fixed broken links to github repos

d6e222f on Nov 1, 2016

3 contributors 

502 lines (495 sloc) 33 KB

Raw

Blame

History







```
1 ## Classification and Regression with H2O Deep Learning
2 #
3 ## Introduction
4 # * Installation and Startup
5 # * Decision Boundaries
6 ## Cover Type Dataset
7 # * Exploratory Data Analysis
8 # * Deep Learning Model
9 # * Hyper-Parameter Search
10 # * Checkpointing
11 # * Cross-Validation
12 # * Model Save & Load
13 ## Regression and Binary Classification
14 ## Deep Learning Tips & Tricks
15 #
16 ### Introduction
17 #This tutorial shows how a H2O [Deep Learning](http://en.wikipedia.org/wiki/Deep_learning) model can be used to do supervised classification
18 #
19 #If run from plain R, execute R in the directory of this script. If run from RStudio, be sure to setwd() to the location of this script. h2o
20 #
21 #More examples and explanations can be found in our [H2O Deep Learning booklet](http://h2o.ai/resources/) and on our [H2O Github Repository]
22 #
23 ##### H2O R Package
24 #
25 #Load the H2O R package:
26 #
27 ## R installation instructions are at http://h2o.ai/download
28 library(h2o)
29 #
30 ##### Start H2O
31 #Start up a 1-node H2O server on your local machine, and allow it to use all CPU cores and up to 2GB of memory:
32 #
33 h2o.init(nthreads=-1, max_mem_size="2G")
34 h2o.removeAll() ## clean slate - just in case the cluster was already running
35 #
36 #The `h2o.deeplearning` function fits H2O's Deep Learning models from within R.
37 #We can run the example from the man page using the `example` function, or run a longer demonstration from the `h2o` package using the `demo`
38 #
39 args(h2o.deeplearning)
40 help(h2o.deeplearning)
41 example(h2o.deeplearning)
42 #demo(h2o.deeplearning) #requires user interaction
43 #
44 #While H2O Deep Learning has many parameters, it was designed to be just as easy to use as the other supervised training methods in H2O. Ea
45 #
46 ##### Let's have some fun first: Decision Boundaries
47 #We start with a small dataset representing red and black dots on a plane, arranged in the shape of two nested spirals. Then we task H2O's
48 #
49 #We visualize the nature of H2O Deep Learning (DL), H2O's tree methods (GBM/DRF) and H2O's generalized linear modeling (GLM) by plotting the
```

```

50 #
51 setwd("~/h2o-tutorials/tutorials/deeplearning") ##For RStudio
52 spiral <- h2o.importFile(path = normalizePath("../data/spiral.csv"))
53 grid <- h2o.importFile(path = normalizePath("../data/grid.csv"))
54 # Define helper to plot contours
55 plotC <- function(name, model, data=spiral, g=grid) {
56   data <- as.data.frame(data) #get data from into R
57   pred <- as.data.frame(h2o.predict(model, g))
58   n=0.5*(sqrt(nrow(g))-1); d <- 1.5; h <- d*(-n:n)/n
59   plot(data[, -3], pch=19, col=data[, 3], cex=0.5,
60        xlim=c(-d,d), ylim=c(-d,d), main=name)
61   contour(h, z=array(ifelse(pred[, 1]=="Red", 0, 1),
62                      dim=c(2*n+1, 2*n+1)), col="blue", lwd=2, add=T)
63 }
64 #
65 #We build a few different models:
66 #
67 #dev.new(noRStudioGD=FALSE) #direct plotting output to a new window
68 par(mfrow=c(2,2)) #set up the canvas for 2x2 plots
69 plotC("DL", h2o.deeplearning(1:2,3,spiral,epochs=1e3))
70 plotC("GBM", h2o.gbm(1:2,3,spiral))
71 plotC("DRF", h2o.randomForest(1:2,3,spiral))
72 plotC("GLM", h2o.glm(1:2,3,spiral,family="binomial"))
73 #
74 #Let's investigate some more Deep Learning models. First, we explore the evolution over training time (number of passes over the data), and
75 #
76 #dev.new(noRStudioGD=FALSE) #direct plotting output to a new window
77 par(mfrow=c(2,2)) #set up the canvas for 2x2 plots
78 ep <- c(1,250,500,750)
79 plotC(paste0("DL ", ep[1], " epochs"),
80       h2o.deeplearning(1:2,3,spiral,epochs=ep[1],
81                        model_id="dl_1"))
82 plotC(paste0("DL ", ep[2], " epochs"),
83       h2o.deeplearning(1:2,3,spiral,epochs=ep[2],
84                        checkpoint="dl_1", model_id="dl_2"))
85 plotC(paste0("DL ", ep[3], " epochs"),
86       h2o.deeplearning(1:2,3,spiral,epochs=ep[3],
87                        checkpoint="dl_2", model_id="dl_3"))
88 plotC(paste0("DL ", ep[4], " epochs"),
89       h2o.deeplearning(1:2,3,spiral,epochs=ep[4],
90                        checkpoint="dl_3", model_id="dl_4"))
91 #
92 #You can see how the network learns the structure of the spirals with enough training time. We explore different network architectures next
93 #
94 #dev.new(noRStudioGD=FALSE) #direct plotting output to a new window
95 par(mfrow=c(2,2)) #set up the canvas for 2x2 plots
96 for (hidden in list(c(11,13,17,19),c(42,42,42),c(200,200),c(1000))) {
97   plotC(paste0("DL hidden=", paste0(hidden, collapse="x")),
98       h2o.deeplearning(1:2,3,spiral,hidden=hidden,epochs=500))
99 }
100 #
101 #It is clear that different configurations can achieve similar performance, and that tuning will be required for optimal performance. Next,
102 #
103 #dev.new(noRStudioGD=FALSE) #direct plotting output to a new window
104 par(mfrow=c(2,2)) #set up the canvas for 2x2 plots
105 for (act in c("Tanh","Maxout","Rectifier","RectifierWithDropout")) {
106   plotC(paste0("DL ", act, " activation"),
107       h2o.deeplearning(1:2,3,spiral,
108                        activation=act, hidden=c(100,100), epochs=1000))
109 }
110 #
111 #Clearly, the dropout rate was too high or the number of epochs was too low for the last configuration, which often ends up performing the
112 #
113 #More information about the parameters can be found in the [H2O Deep Learning booklet](http://h2o.ai/resources/).
114 #
115 ### Cover Type Dataset
116 #We important the full cover type dataset (581k rows, 13 columns, 10 numerical, 3 categorical).

```

```

117 #We also split the data 3 ways: 60% for training, 20% for validation (hyper parameter tuning) and 20% for final testing.
118 #
119 df <- h2o.importFile(path = normalizePath("../data/covtype.full.csv"))
120 dim(df)
121 df
122 splits <- h2o.splitFrame(df, c(0.6,0.2), seed=1234)
123 train <- h2o.assign(splits[[1]], "train.hex") # 60%
124 valid <- h2o.assign(splits[[2]], "valid.hex") # 20%
125 test <- h2o.assign(splits[[3]], "test.hex") # 20%
126 #
127 #Here's a scalable way to do scatter plots via binning (works for categorical and numeric columns) to get more familiar with the dataset.
128 #
129 #dev.new(noRStudioGD=FALSE) #direct plotting output to a new window
130 par(mfrow=c(1,1)) # reset canvas
131 plot(h2o.tabulate(df, "Elevation", "Cover_Type"))
132 plot(h2o.tabulate(df, "Horizontal_Distance_To_Roadways", "Cover_Type"))
133 plot(h2o.tabulate(df, "Soil_Type", "Cover_Type"))
134 plot(h2o.tabulate(df, "Horizontal_Distance_To_Roadways", "Elevation" ))
135 #
136 ##### First Run of H2O Deep Learning
137 #Let's run our first Deep Learning model on the covtype dataset.
138 #We want to predict the `Cover_Type` column, a categorical feature with 7 levels, and the Deep Learning model will be tasked to perform (mu
139 #
140 response <- "Cover_Type"
141 predictors <- setdiff(names(df), response)
142 predictors
143 #
144 #To keep it fast, we only run for one epoch (one pass over the training data).
145 #
146 m1 <- h2o.deeplearning(
147   model_id="dl_model_first",
148   training_frame=train,
149   validation_frame=valid, ## validation dataset: used for scoring and early stopping
150   x=predictors,
151   y=response,
152   #activation="Rectifier", ## default
153   #hidden=c(200,200), ## default: 2 hidden layers with 200 neurons each
154   epochs=1,
155   variable_importances=T ## not enabled by default
156 )
157 summary(m1)
158 #
159 #Inspect the model in [Flow](http://localhost:54321/) for more information about model building etc. by issuing a cell with the content `ge
160 #
161 ##### Variable Importances
162 #Variable importances for Neural Network models are notoriously difficult to compute, and there are many [pitfalls](ftp://ftp.sas.com/pub/n
163 #
164 head(as.data.frame(h2o.varimp(m1)))
165 #
166 ##### Early Stopping
167 #Now we run another, smaller network, and we let it stop automatically once the misclassification rate converges (specifically, if the movi
168 #
169 m2 <- h2o.deeplearning(
170   model_id="dl_model_faster",
171   training_frame=train,
172   validation_frame=valid,
173   x=predictors,
174   y=response,
175   hidden=c(32,32,32), ## small network, runs faster
176   epochs=100000, ## hopefully converges earlier...
177   score_validation_samples=10000, ## sample the validation dataset (faster)
178   stopping_rounds=2,
179   stopping_metric="misclassification", ## could be "MSE", "logloss", "r2"
180   stopping_tolerance=0.01
181 )
182 summary(m2)
183 plot(m2)

```

```

184 #
185 ##### Adaptive Learning Rate
186 #By default, H2O Deep Learning uses an adaptive learning rate ([ADADELTA](http://arxiv.org/pdf/1212.5701v1.pdf)) for its stochastic gradient
187 #
188 #If `adaptive_rate` is disabled, several manual learning rate parameters become important: `rate`, `rate_annealing`, `rate_decay`, `momentum`
189 #
190 ##### Tuning
191 #With some tuning, it is possible to obtain less than 10% test set error rate in about one minute. Error rates of below 5% are possible with
192 #
193 m3 <- h2o.deeplearning(
194   model_id="dl_model_tuned",
195   training_frame=train,
196   validation_frame=valid,
197   x=predictors,
198   y=response,
199   overwrite_with_best_model=F, ## Return the final model after 10 epochs, even if not the best
200   hidden=c(128,128,128),      ## more hidden layers -> more complex interactions
201   epochs=10,                  ## to keep it short enough
202   score_validation_samples=10000, ## downsample validation set for faster scoring
203   score_duty_cycle=0.025,      ## don't score more than 2.5% of the wall time
204   adaptive_rate=F,             ## manually tuned learning rate
205   rate=0.01,
206   rate_annealing=2e-6,
207   momentum_start=0.2,         ## manually tuned momentum
208   momentum_stable=0.4,
209   momentum_ramp=1e7,
210   l1=1e-5,                    ## add some L1/L2 regularization
211   l2=1e-5,
212   max_w2=10                    ## helps stability for Rectifier
213 )
214 summary(m3)
215 #
216 #Let's compare the training error with the validation and test set errors
217 #
218 h2o.performance(m3, train=T)      ## sampled training data (from model building)
219 h2o.performance(m3, valid=T)      ## sampled validation data (from model building)
220 h2o.performance(m3, newdata=train) ## full training data
221 h2o.performance(m3, newdata=valid) ## full validation data
222 h2o.performance(m3, newdata=test)  ## full test data
223 #
224 #To confirm that the reported confusion matrix on the validation set (here, the test set) was correct, we make a prediction on the test set
225 #
226 pred <- h2o.predict(m3, test)
227 pred
228 test$Accuracy <- pred$predict == test$Cover_Type
229 1-mean(test$Accuracy)
230 #
231 ##### Hyper-parameter Tuning with Grid Search
232 #Since there are a lot of parameters that can impact model accuracy, hyper-parameter tuning is especially important for Deep Learning:
233 #
234 #For speed, we will only train on the first 10,000 rows of the training dataset:
235 #
236 sampled_train=train[1:10000,]
237 #
238 #The simplest hyperparameter search method is a brute-force scan of the full Cartesian product of all combinations specified by a grid search
239 #
240 hyper_params <- list(
241   hidden=list(c(32,32,32),c(64,64)),
242   input_dropout_ratio=c(0,0.05),
243   rate=c(0.01,0.02),
244   rate_annealing=c(1e-8,1e-7,1e-6)
245 )
246 hyper_params
247 grid <- h2o.grid(
248   algorithm="deeplearning",
249   grid_id="dl_grid",
250   training_frame=sampled_train,

```

```

251 validation_frame=valid,
252 x=predictors,
253 y=response,
254 epochs=10,
255 stopping_metric="misclassification",
256 stopping_tolerance=1e-2,      ## stop when misclassification does not improve by >=1% for 2 scoring events
257 stopping_rounds=2,
258 score_validation_samples=10000, ## downsample validation set for faster scoring
259 score_duty_cycle=0.025,      ## don't score more than 2.5% of the wall time
260 adaptive_rate=F,            ## manually tuned learning rate
261 momentum_start=0.5,         ## manually tuned momentum
262 momentum_stable=0.9,
263 momentum_ramp=1e7,
264 l1=1e-5,
265 l2=1e-5,
266 activation=c("Rectifier"),
267 max_w2=10,                  ## can help improve stability for Rectifier
268 hyper_params=hyper_params
269 )
270 grid
271 #
272 #Let's see which model had the lowest validation error:
273 #
274 grid <- h2o.getGrid("dl_grid",sort_by="err",decreasing=FALSE)
275 grid
276
277 ## To see what other "sort_by" criteria are allowed
278 #grid <- h2o.getGrid("dl_grid",sort_by="wrong_thing",decreasing=FALSE)
279
280 ## Sort by logloss
281 h2o.getGrid("dl_grid",sort_by="logloss",decreasing=FALSE)
282
283 ## Find the best model and its full set of parameters
284 grid@summary_table[1,]
285 best_model <- h2o.getModel(grid@model_ids[[1]])
286 best_model
287
288 print(best_model@allparameters)
289 print(h2o.performance(best_model, valid=T))
290 print(h2o.logloss(best_model, valid=T))
291 #
292 ##### Random Hyper-Parameter Search
293 #Often, hyper-parameter search for more than 4 parameters can be done more efficiently with random parameter search than with grid search.
294 #
295 hyper_params <- list(
296   activation=c("Rectifier","Tanh","Maxout","RectifierWithDropout","TanhWithDropout","MaxoutWithDropout"),
297   hidden=list(c(20,20),c(50,50),c(30,30,30),c(25,25,25,25)),
298   input_dropout_ratio=c(0,0.05),
299   l1=seq(0,1e-4,1e-6),
300   l2=seq(0,1e-4,1e-6)
301 )
302 hyper_params
303
304 ## Stop once the top 5 models are within 1% of each other (i.e., the windowed average varies less than 1%)
305 search_criteria = list(strategy = "RandomDiscrete", max_runtime_secs = 360, max_models = 100, seed=1234567, stopping_rounds=5, stopping_tol
306 dl_random_grid <- h2o.grid(
307   algorithm="deeplearning",
308   grid_id = "dl_grid_random",
309   training_frame=sampled_train,
310   validation_frame=valid,
311   x=predictors,
312   y=response,
313   epochs=1,
314   stopping_metric="logloss",
315   stopping_tolerance=1e-2,      ## stop when logloss does not improve by >=1% for 2 scoring events
316   stopping_rounds=2,
317   score_validation_samples=10000, ## downsample validation set for faster scoring

```

```

318   score_duty_cycle=0.025,          ## don't score more than 2.5% of the wall time
319   max_w2=10,                      ## can help improve stability for Rectifier
320   hyper_params = hyper_params,
321   search_criteria = search_criteria
322 )
323 grid <- h2o.getGrid("dl_grid_random",sort_by="logloss",decreasing=FALSE)
324 grid
325
326 grid@summary_table[1,]
327 best_model <- h2o.getModel(grid@model_ids[[1]]) ## model with lowest logloss
328 best_model
329 #
330 #Let's look at the model with the lowest validation misclassification rate:
331 #
332 grid <- h2o.getGrid("dl_grid",sort_by="err",decreasing=FALSE)
333 best_model <- h2o.getModel(grid@model_ids[[1]]) ## model with lowest classification error (on validation, since it was available during tra
334 h2o.confusionMatrix(best_model,valid=T)
335 best_params <- best_model@allparameters
336 best_params$activation
337 best_params$hidden
338 best_params$input_dropout_ratio
339 best_params$l1
340 best_params$l2
341 #
342 #####Checkpointing
343 #Let's continue training the manually tuned model from before, for 2 more epochs. Note that since many important parameters such as `epochs
344 #
345 max_epochs <- 12 ## Add two more epochs
346 m_cont <- h2o.deeplearning(
347   model_id="dl_model_tuned_continued",
348   checkpoint="dl_model_tuned",
349   training_frame=train,
350   validation_frame=valid,
351   x=predictors,
352   y=response,
353   hidden=c(128,128,128),          ## more hidden layers -> more complex interactions
354   epochs=max_epochs,              ## hopefully long enough to converge (otherwise restart again)
355   stopping_metric="logloss",      ## logloss is directly optimized by Deep Learning
356   stopping_tolerance=1e-2,        ## stop when validation logloss does not improve by >=1% for 2 scoring events
357   stopping_rounds=2,
358   score_validation_samples=10000, ## downsample validation set for faster scoring
359   score_duty_cycle=0.025,         ## don't score more than 2.5% of the wall time
360   adaptive_rate=F,               ## manually tuned learning rate
361   rate=0.01,
362   rate_annealing=2e-6,
363   momentum_start=0.2,            ## manually tuned momentum
364   momentum_stable=0.4,
365   momentum_ramp=1e7,
366   l1=1e-5,                      ## add some L1/L2 regularization
367   l2=1e-5,
368   max_w2=10                      ## helps stability for Rectifier
369 )
370 summary(m_cont)
371 plot(m_cont)
372 #
373 #Once we are satisfied with the results, we can save the model to disk (on the cluster). In this example, we store the model in a directory
374 #
375 path <- h2o.saveModel(m_cont,
376   path="./mybest_deeplearning_covtype_model", force=TRUE)
377 #
378 #It can be loaded later with the following command:
379 #
380 print(path)
381 m_loaded <- h2o.loadModel(path)
382 summary(m_loaded)
383 #
384 #This model is fully functional and can be inspected, restarted, or used to score a dataset, etc. Note that binary compatibility between H2

```

```

385 #
386 #####Cross-Validation
387 #For N-fold cross-validation, specify `nfolds>1` instead of (or in addition to) a validation frame, and `N+1` models will be built: 1 model
388 #
389 dlmodel <- h2o.deeplearning(
390   x=predictors,
391   y=response,
392   training_frame=train,
393   hidden=c(10,10),
394   epochs=1,
395   nfolds=5,
396   fold_assignment="Modulo" # can be "AUTO", "Modulo", "Random" or "Stratified"
397 )
398 dlmodel
399 #
400 #N-fold cross-validation is especially useful with early stopping, as the main model will pick the ideal number of epochs from the converge
401 #
402 ###Regression and Binary Classification
403 #Assume we want to turn the multi-class problem above into a binary classification problem. We create a binary response as follows:
404 #
405 train$bin_response <- ifelse(train[,response]=="class_1", 0, 1)
406 #
407 #Let's build a quick model and inspect the model:
408 #
409 dlmodel <- h2o.deeplearning(
410   x=predictors,
411   y="bin_response",
412   training_frame=train,
413   hidden=c(10,10),
414   epochs=0.1
415 )
416 summary(dlmodel)
417 #
418 #Instead of a binary classification model, we find a regression model (`H2ORegressionModel`) that contains only 1 output neuron (instead of
419 #H2O Deep Learning supports regression for distributions other than `Gaussian` such as `Poisson`, `Gamma`, `Tweedie`, `Laplace`. It also su
420 #
421 #To perform classification, the response must first be turned into a categorical (factor) feature:
422 #
423 train$bin_response <- as.factor(train$bin_response) ##make categorical
424 dlmodel <- h2o.deeplearning(
425   x=predictors,
426   y="bin_response",
427   training_frame=train,
428   hidden=c(10,10),
429   epochs=0.1
430   #balance_classes=T    ## enable this for high class imbalance
431 )
432 summary(dlmodel) ## Now the model metrics contain AUC for binary classification
433 plot(h2o.performance(dlmodel)) ## display ROC curve
434 #
435 #Now the model performs (binary) classification, and has multiple (2) output neurons.
436 #
437 ###Unsupervised Anomaly detection
438 #For instructions on how to build unsupervised models with H2O Deep Learning, we refer to our previous [Tutorial on Anomaly Detection with
439 #[Unsupervised Pretraining with an AutoEncoder R code example](https://github.com/h2oai/h2o-3/blob/master/h2o-r/tests/testdir_algos/deeplea
440 #
441 #
442 ###H2O Deep Learning Tips & Tricks
443 #
444 #####Performance Tuning
445 #The [Definitive H2O Deep Learning Performance Tuning](http://blog.h2o.ai/2015/08/deep-learning-performance-august/) blog post covers many
446 #
447 #####Activation Functions
448 #While sigmoids have been used historically for neural networks, H2O Deep Learning implements `Tanh`, a scaled and shifted variant of the s
449 #
450 #####Generalization Techniques
451 #L1 and L2 penalties can be applied by specifying the `l1` and `l2` parameters. Intuition: L1 lets only strong weights survive (constant pu

```

```
452 #
453 #####Early stopping and optimizing for lowest validation error
454 #By default, Deep Learning training stops when the `stopping_metric` does not improve by at least `stopping_tolerance` (0.01 means 1% impro
455 #
456 #####Training Samples per MapReduce Iteration
457 #The parameter `train_samples_per_iteration` matters especially in multi-node operation. It controls the number of rows trained on for each
458 #
459 #####Categorical Data
460 #For categorical data, a feature with K factor levels is automatically one-hot encoded (horizontalized) into K-1 input neurons. Hence, the
461 #
462 #####Sparse Data
463 #If the input data is sparse (many zeros), then it might make sense to enable the `sparse` option. This will result in the input not being
464 #
465 #####Missing Values
466 #H2O Deep Learning automatically does mean imputation for missing values during training (leaving the input layer activation at 0 after sta
467 #
468 #####Loss functions, Distributions, Offsets, Observation Weights
469 #H2O Deep Learning supports advanced statistical features such as multiple loss functions, non-Gaussian distributions, per-row offsets and
470 #In addition to `Gaussian` distributions and `Squared` loss, H2O Deep Learning supports `Poisson`, `Gamma`, `Tweedie` and `Laplace` distrib
471 #
472 #We refer to our [H2O Deep Learning R test code examples](https://github.com/h2oai/h2o-3/tree/master/h2o-r/tests/testdir_algos/deeplearning
473 #
474 #####Exporting Weights and Biases
475 #The model parameters (weights connecting two adjacent layers and per-neuron bias terms) can be stored as H2O Frames (like a dataset) by en
476 #
477 iris_dl <- h2o.deeplearning(1:4,5,as.h2o(iris),
478                             export_weights_and_biases=T)
479 h2o.weights(iris_dl, matrix_id=1)
480 h2o.weights(iris_dl, matrix_id=2)
481 h2o.weights(iris_dl, matrix_id=3)
482 h2o.biases(iris_dl, vector_id=1)
483 h2o.biases(iris_dl, vector_id=2)
484 h2o.biases(iris_dl, vector_id=3)
485 #plot weights connecting `Sepal.Length` to first hidden neurons
486 plot(as.data.frame(h2o.weights(iris_dl, matrix_id=1))[,1])
487 #
488 #####Reproducibility
489 #Every run of DeepLearning results in different results since multithreading is done via [Hogwild!](http://www.eecs.berkeley.edu/~brecht/pa
490 #
491 #####Scoring on Training/Validation Sets During Training
492 #The training and/or validation set errors *can* be based on a subset of the training or validation data, depending on the values for `score
493 #
494 #Note that the default value of `score_duty_cycle=0.1` limits the amount of time spent in scoring to 10%, so a large number of scoring samp
495 #
496 #Stratified sampling of the validation dataset can help with scoring on datasets with class imbalance. Note that this option also requires
497 #
498 ##### More information can be found in the [H2O Deep Learning booklet](http://h2o.ai/resources/), in our [H2O SlideShare Presentations](http
499 #
500 ##### All done, shutdown H2O
501 h2o.shutdown(prompt=FALSE)
```

